

# M N I M E

## Multimodal Neural Interface Machine Extension

*nigh.mh*

*“... explaining how the Galaxy could be colonised, not by a vast fleet of starships, but by launching into deep space one single intelligent machine.”*

— Adrian Berry, *The Super-Intelligent Machine* (Berry, 1983)

Kyle Bauer

kyle.bauer@bellevuecollege.edu

### Abstract

Document processing software has long relied on cloud-based pipelines or heavyweight external dependencies, raising concerns around data privacy, latency, and offline operability. We present **MNIME** (Multimodal Neural Interface Machine Extension), a fully local, offline desktop system that unifies high-throughput PDF and image processing with an integrated, fine-tuned large language model (LLM) running entirely on-device. Built with Python and PyQt6, MNIME combines a C-accelerated PDF engine (PyMuPDF), a FAISS-backed semantic retrieval layer with persistent global memory caching, and MNIME-Core-V5-Q4\_K\_M, a quantized 1.5B-parameter model fine-tuned from Qwen2.5-1.5B-Instruct featuring dynamic LoRA adapter injection for document-centric reasoning tasks including cross-referencing, semantic bookmarking, and conversational retrieval-augmented generation (RAG). We describe the system architecture, fine-tuning methodology, multimodal interface design, and performance characteristics across ten operational modes. On-device empirical telemetry demonstrates that MNIME achieves sub-second UI responsiveness and sustained token generation throughput of 56.15 tokens/s on consumer hardware, while theoretical projections indicate competitive document comprehension and cross-reference quality, transmitting zero bytes of user data externally. Our results highlight the viability of fully local multimodal document AI for privacy-sensitive institutional and personal use cases.

## 1 Introduction

Modern document workflows increasingly demand intelligent interaction: users want to query, summarize, and cross-reference PDFs without leaving a single interface. Existing solutions typically route document content through cloud inference end-

points (OpenAI, 2024; Gemini Team et al., 2025), creating hard barriers for users operating under data-sovereignty constraints, network-limited environments, or simply preferring that sensitive documents never leave their device.

Locally-hosted LLMs have become practical for consumer hardware following the quantization advances of GGUF/llama.cpp (Gerganov, 2023) and the release of sub-2B-parameter instruction-tuned models such as Qwen2.5-1.5B-Instruct (Qwen Team, 2025). However, integrating a capable local LLM into a cohesive, multimodal desktop application—one that handles PDF merging, editing, format conversion, semantic search, and cross-referencing—remains largely unexplored in publicly documented systems.

This paper introduces **MNIME**, which addresses this gap through three core contributions:

1. A **ten-mode multimodal processing pipeline** covering PDF merge, edit, compress, convert, and intelligent bookmark generation, executed via native C-level routines for minimal latency.
2. **MNIME-Core-V5**, a fine-tuned GGUF model specialized for document reasoning, cross-referencing, and retrieval-augmented conversational question answering over user-loaded corpora.
3. A **premium local-first UI** built with PyQt6 featuring physics-based particle transitions, a magnetic translucent pop-out chat, 4.0× Retina-quality document rendering, and an O(1) drag-and-drop carousel—all with zero network dependency.

Figure 1 illustrates the high-level architecture. Section 2 reviews related work. Section 3 details

the system design. Section 4 describes the MNIME-Core-V5 fine-tuning. Section 6 presents projected performance characteristics. Section 8 concludes.

## 2 Related Work

### 2.1 Local Inference and Quantized LLMs

The proliferation of 4-bit quantized GGUF models via llama.cpp (Gerganov, 2023) has made sub-8GB-VRAM inference of billion-parameter models feasible. GPT4All (Anand et al., 2023) and LM Studio popularized local chat interfaces but lack document-processing pipelines. LlamaIndex (Liu, 2022) and LangChain (Chase, 2022) provide RAG orchestration but require Python runtime configuration and are not designed as end-user desktop applications.

### 2.2 Document AI and PDF Processing

PyMuPDF (Artifex Software, 2024) provides a C-accelerated binding to the MuPDF engine, enabling lossless PDF manipulation at speeds unattainable with pure-Python libraries such as PyPDF2. Semantic PDF understanding has been explored in LayoutLMv3 (Huang et al., 2022), but these systems require GPU clusters for training and cloud APIs for deployment.

### 2.3 Desktop Document Assistants

Commercial tools such as Adobe Acrobat AI Assistant and Notion AI provide document summarization but depend on external servers. Open-source alternatives such as Ollama provide local inference but lack native PDF manipulation. MNIME uniquely combines both dimensions: native document processing and integrated local neural inference within a single installable application.

### 2.4 Retrieval-Augmented Generation

RAG (Lewis et al., 2020) enables LLMs to answer questions grounded in retrieved passages, reducing hallucination. FAISS (Johnson et al., 2019) provides efficient vector similarity search for large corpora. We integrate both into MNIME’s NLP engine, constructing per-session FAISS indices from document text chunks encoded by a local sentence embedding model.

## 3 System Architecture

MNIME is composed of three primary layers: the **Core Processing Engine**, the **NLP/RAG Engine**,

and the **UI Layer**. The modular architecture ensures each layer can operate independently, enabling lightweight non-NLP usage with instant startup.

### 3.1 Core Processing Engine

The processing engine (`core/pdf_engine.py`) implements ten operational modes via PyMuPDF C-level routines:

1. **Merge**: High-throughput batch merge of up to 5,000 PDF/image files with chunked segment memory management and optional NLP-driven outline/TOC generation.
2. **JPG → PDF**: Lossless image-to-PDF conversion.
3. **TEXT → PDF**: Native vector text PDF generation.
4. **PDF → JPG**: High-resolution page rasterization.
5. **Split PDF**: Page-level decomposition with optional NLP-powered smart naming.
6. **Compress**: Content stream deduplication and object compression.
7. **PDF → DOCX**: Layout-preserving Word export.
8. **Semantic Bookmark**: Typography-driven chapter detection with NLP-generated summaries.
9. **NLP/RAG Chat**: Conversational query over loaded documents.
10. **Cross-Reference**: Highlighted-section synthesis briefs against multiple documents.

To prevent system memory and handle exhaustion during high-volume document aggregation (e.g., merging 5,000 multi-page PDFs totaling over 100,000 pages), the engine employs a *chunked segment merge architecture*. Batches exceeding 250 documents are split into temporary intermediate PDF segments on disk, releasing C-level memory allocations before executing a lightweight final assembly pass. Furthermore, when NLP indexing is active, NLP Engine inspects document text during merge passes to dynamically synthesize concise AI section titles (`generate_smart_filename`) and embed a unified, navigable Table of Contents

(TOC) into the generated document’s outline meta-data.

All processing runs in QThread workers (core/worker.py), maintaining 60 FPS UI responsiveness with real-time progress indicators.

### 3.2 NLP and RAG Engine

The NLP engine (core/nlp\_engine.py) loads MNIME-Core-V5-Q4\_K\_M.gguf via llama-cpp-python with automatic NVIDIA GPU layer offloading (n\_gpu\_layers=-1). On systems without a GPU, the model runs fully in CPU mode with acceptable latency for conversational interaction.

The semantic search engine (core/search\_engine.py) extracts text chunks from all open PDFs, encodes them using a local sentence embedding model (BAAI/bge-small-en-v1.5), and indexes them in a FAISS flat-index. Crucially, these vectors are automatically appended to a persistent global vector store in the user’s local application data directory. At query time, up to two additional passages from this store (excluding duplicates of the open-document results) are added to the context, enabling recall of previously ingested documents across sessions. At inference time, the user query is embedded and the top- $k$  retrieved chunks are prepended to the model’s context window.

**Memory management:** The engine optionally loads a user-selected LoRA adapter (GGUF) on top of the base model; selecting or removing an adapter reloads the model on a background thread. The model and local session FAISS index are held exclusively in memory during active NLP sessions and are completely deallocated upon toggle-off or application exit, preventing background memory consumption.

### 3.3 UI Layer

The UI (ui/main\_window.py) is a frameless, free-floating dark metallic window with no system chrome. Key UI components are listed in Table 1.

The carousel supports  $O(1)$  drag-and-drop re-ordering through surgical layout index shifts, avoiding widget teardown/rebuild for even 5,000 loaded files. Thumbnails and SVG icons are rasterized once and cached, eliminating CPU resampling during scroll events.

| Component          | Description                            |
|--------------------|----------------------------------------|
| carousel_view.py   | $O(1)$ horizontal file card carousel   |
| file_card.py       | Thumbnail, status badge, drag-and-drop |
| nlp_view.py        | RAG conversational interface           |
| document_viewer.py | 4.0× Retina document viewer            |
| merge_particles.py | Physics particle simulation            |
| reader_dialog.py   | Independent frameless reader window    |
| file_dialog.py     | Custom dark-mode file explorer         |
| nerds.py           | Real-time telemetry & performance HUD  |

Table 1: Key UI components of MNIME.

## 4 MNIME-Core-V5 Fine-Tuning

### 4.1 Base Model and Quantization

We select **Qwen2.5-1.5B-Instruct** (Qwen Team, 2025) as the base model due to its strong instruction-following performance at the 1.5B parameter scale, its permissive license, and its compatibility with llama.cpp’s GGUF quantization pipeline. Post fine-tuning, we quantize to Q4\_K\_M, which provides an optimal balance between model quality and memory footprint ( $\approx 1.0$  GB).

### 4.2 Training Data and Multi-Stage Alignment Pipeline

To ensure MNIME-Core-V5 remains highly responsive, accurate, and empathetic, we structured the fine-tuning process across a multi-stage, hybrid compute pipeline. This separates foundational QA alignment (trained locally) from expansive philosophical, ethical, and anti-bias instruction (trained via high-compute cloud instances and fused locally).

#### Stage 1: Foundational V3 Local Training.

The foundational V3 phase aligns the base Qwen2.5-1.5B-Instruct weights to MNIME’s core identity, focusing on closed QA, document summarization, and information extraction. Using a highly filtered 5,482-example subset of Databricks Dolly 15k (training/mnime\_v3\_dataset\_clean.jsonl), the model was trained locally on an NVIDIA RTX 3060 (12 GB VRAM). The pipeline employs HuggingFace trl, peft, and bitsandbytes to execute 4-bit NF4 Low-Rank Adaptation (LoRA) (Hu et al., 2021).

#### Stage 2: V4 Philosophy & Robustness Cloud Upgrade.

Following the foundational local training, the model weights were elevated through a cloud fine-tuning phase on Google Cloud Platform (GCP) accelerators. The

V4 phase utilized a synthetically engineered 15,000-example dataset (training/mnime\_v4\_philosophy\_dataset\_clean.jsonl) designed to impart ontological discourse (Stoicism, Existentialism) and conversational empathy. Concurrently, adversarial prompt injections and hallucination traps were embedded directly into the training corpus to hardcode resilience against jailbreaks. The resulting adapter was merged back into the base matrix in full fp16 precision.

**Stage 3: V5 Ethics, Empathy, & Anti-Bias Alignment.** The culminating alignment phase directly eliminates prejudice, stereotypes, and systemic bigotry. Using a 20,000-example synthetic dataset (training/mnime\_v5\_antibias\_dataset\_clean.jsonl), MNIME-Core-V5 was trained to actively deconstruct prejudiced statements using sociological evidence and empirical logic rather than relying on evasive, generic AI refusal templates. All training examples were algorithmically scrubbed of superficial chatbot phraseology to ensure the final MNIME persona remains grounded, authentically humanistic, and intellectually rigorous.

### 4.3 The Neural Assimilation Engine

To ensure MNIME is not merely a static interface but an actively evolving cognitive architecture, the application integrates a *Neural Assimilation Engine* (NAE). While MNIME-Core-V5 operates as a resilient logical scribe, users can asynchronously digest and fuse external, specialized domain weights into the core matrix.

Operating directly on full-precision (fp16/bf16) HuggingFace safetensors representations, the NAE utilizes streaming tensor-level fusion techniques (such as TIES-merging and linear task-vector interpolation). The engine isolates parameter deltas ( $\Delta W = W_{\text{donor}} - W_{\text{base}}$ ), applies magnitude density pruning, resolves sign conflicts through majority voting, and assimilates the high-value specialized weights without overwriting core reasoning pathways or inducing catastrophic forgetting. Once fused, the resulting matrix is converted directly to multi-precision GGUF targets (including Q8\_0 and Q4\_K\_M) for immediate local execution.

Key hyperparameters and system specifications for the initial local training phase are summarized in Table 2.

| Parameter / Setting  | Value                                       |
|----------------------|---------------------------------------------|
| Base model           | Qwen2.5-1.5B-Instruct                       |
| Hardware             | AMD Ryzen 9 7950X,<br>RTX 3060 Ventus 12 GB |
| System memory        | 64 GB DDR5 6000 MHz                         |
| Framework            | unsloth (LoRA)                              |
| LoRA rank / $\alpha$ | $r = 32, \alpha = 32$                       |
| LoRA targets         | attention + MLP projections                 |
| Training examples    | 5,482                                       |
| Learning rate        | $2 \times 10^{-4}$ (linear, 100 warmup)     |
| Batch / Grad accum   | 4 / 4                                       |
| Training steps       | 1 epoch ( $\approx 1,370$ steps)            |
| Sequence length      | 2,048 tokens                                |
| Inference context    | 4,096 tokens                                |
| Quantization         | Q4_K_M (post-training)                      |

Table 2: MNIME-Core-V5 training and deployment configuration.

## 5 Architecture




Figure 1: High-level MNIME architecture. The PyQt6 UI dispatches to either the C-accelerated PDF engine (left) or the local NLP engine (right). All inference and document processing occur on-device with no external network calls.

## 6 Performance Characteristics

### 6.1 Projection Basis and Methodology

*The operational speed and quality figures presented in Sections 6.2–6.4 are design-informed, theoretical projections rather than laboratory-benchmarked results. They are derived from the algorithmic complexity of the implemented data structures ( $O(1)$  carousel operations, C-level PyMuPDF routines), published throughput benchmarks for the underlying libraries (PyMuPDF, llama-cpp-python, FAISS), and typical generation speeds reported for Q4\_K\_M quantized 1.5B-parameter GGUF models on the specified reference hardware. NLP quality estimates extrapolate from published LoRA fine-tuning scaling results at comparable rank and step counts. Hardware latency, memory consumption, and prompt ingestion throughput were subsequently empirically measured and validated on-device as detailed in Section 6.5.*

All projections assume the reference workstation on which the system was developed, trained, and

tested: an AMD Ryzen 9 7950X 16-Core processor, 64 GB DDR5 6000 MHz RAM, and an NVIDIA GeForce RTX 3060 Ventus 12 GB GPU.

## 6.2 UI Responsiveness (Projected)

Table 3 reports projected frame times for carousel operations at varying file counts. Because reordering is implemented as a surgical layout-index shift with no widget teardown, time complexity is  $O(1)$  in file count; frame times are expected to remain well below the 16.7 ms threshold for 60 FPS regardless of corpus size.

| File Count | Reorder (ms) | Scroll (ms) |
|------------|--------------|-------------|
| 100        | 0.8          | 1.2         |
| 500        | 1.1          | 1.5         |
| 1,000      | 1.4          | 2.0         |
| 2,500      | 2.3          | 2.8         |
| 5,000      | 4.1          | 4.9         |

Table 3: (*Projected*) UI frame time (ms) for carousel operations on the reference hardware.  $O(1)$  complexity ensures times stay well below the 16.7 ms bound for 60 FPS at all file counts.

## 6.3 Document Processing Throughput (Projected)

Table 4 shows projected throughput for each processing mode applied to a representative 100-page, 15 MB test PDF. Figures are derived from published PyMuPDF C-level operation speeds scaled to the reference document size.

| Mode                   | Time (s) | Output     |
|------------------------|----------|------------|
| PDF Merge (10 files)   | 0.41     | 148 MB     |
| Compress (100 pp.)     | 1.23     | 9.2 MB     |
| PDF $\rightarrow$ JPG  | 2.81     | 100 images |
| JPG $\rightarrow$ PDF  | 0.19     | 4.3 MB     |
| PDF $\rightarrow$ DOCX | 3.64     | 1.1 MB     |
| Semantic Bookmark      | 5.12     | 18 entries |

Table 4: (*Projected*) Processing throughput per operational mode on the reference hardware.

## 6.4 NLP Task Quality (Projected)

Table 5 provides projected NLP quality scores for cross-referencing (ROUGE-L) and document-grounded QA (F1). The GPT-4o-mini figures are drawn from published benchmarks (OpenAI, 2024). The Qwen2.5-1.5B base model figures reflect published instruction-following results for that model family. The MNIME-Core projection extrapolates the expected gain from LoRA fine-tuning at rank 16 over 2,000 steps on domain-specific data, following published LoRA scaling results (Hu et al., 2021).

| Model                        | XRef<br>(ROUGE-L) | DocQA<br>(F1) | Local      |
|------------------------------|-------------------|---------------|------------|
| GPT-4o-mini*                 | 0.61              | 0.79          | No         |
| Qwen2.5-1.5B (base)*         | 0.38              | 0.51          | Yes        |
| <b>MNIME-Core-V5</b> (proj.) | <b>0.54</b>       | <b>0.71</b>   | <b>Yes</b> |

Table 5: (*Projected*) Theoretical NLP quality on cross-referencing (XRef) and document-grounded QA (DocQA). \*Published benchmark figures; MNIME-Core-V5 values are theoretical projections extrapolated from LoRA fine-tuning scaling results.

Projected MNIME-Core-V5 performance represents  $\approx 88\%$  of GPT-4o-mini’s cross-referencing quality and  $\approx 90\%$  of its document QA F1, with zero network dependency. The expected improvement over the base model is grounded in published results showing consistent gains from domain-specific LoRA fine-tuning at comparable parameter counts and step budgets.

## 6.5 Memory and Latency (Projected and Empirical)

In lightweight mode (NLP disabled), overall process memory footprint stabilizes at  $\approx 1.0$  GB, encompassing the PyQt6 desktop framework and C-accelerated document processing engines. When MNIME-Core-V5 is activated, total process working set increases to  $\approx 1.8$ – $1.9$  GB, with  $\approx 800$ – $900$  MB dedicated to quantized weights and tensor contexts. Based on published llama-cpp-python throughput for Q4\_K\_M 1.5B models, first-token latency was initially projected at  $\approx 0.38$  s (GPU) and  $\approx 2.4$  s (CPU-only), with sustained generation at  $\approx 48$  tokens/s (GPU).

**Empirical Validation.** To validate these projections, direct on-device empirical benchmarks were conducted on the reference workstation (AMD Ryzen 9 7950X, 64 GB DDR5 6000 MHz RAM, NVIDIA GeForce RTX 3060 Ventus 12 GB) using the bundled MNIME-Core-V5 model. The measured empirical performance validated the feasibility of sub-2B local AI: the model achieved a prompt-context evaluation throughput of **495.89 tokens/s** (151.24 ms for 75 tokens), a first-token latency of **0.08 s**, and a sustained generation throughput of **56.15 tokens/s** (17.81 ms per token), utilizing a 302.75 MB compute buffer. Furthermore, end-to-end integration tests confirmed sub-second heterogeneous document composition (vector PDF, multi-page paginated text, and rasterized images merged in 0.41 s) with 100% chapter-level TOC heuristic extraction precision. Addition-

ally, a high-volume stability test demonstrated that the two-tier chunked segment architecture successfully merged 5,000 PDFs without memory exhaustion. The system processed 250-document chunks in approximately 21 seconds per chunk, maintaining a stable active RAM footprint while concurrently executing NLP-driven TOC generation. Users can reproduce and externally validate these metrics directly via the bundled `benchmark.py` GUI utility.

## 7 Discussion

**Privacy and Data Sovereignty.** MNIME’s fully local architecture is a deliberate design choice. In healthcare, legal, and education settings, documents may contain protected information that cannot legally or ethically be transmitted to third-party cloud services. MNIME provides these users with modern AI document assistance without requiring any privacy trade-offs.

**Fine-Tuning Payoff.** The projected 16-point ROUGE-L improvement of MNIME-Core-V5 over the Qwen2.5-1.5B base model on cross-referencing (Table 5) is consistent with published LoRA fine-tuning results at comparable rank and step budgets (Hu et al., 2021), suggesting that domain-specific fine-tuning—even at modest scale—is likely to yield meaningful gains for specialized tasks. This aligns with findings in instruction tuning literature (Wei et al., 2022; Ouyang et al., 2022).

**Limitations.** MNIME-Core-V5’s 1.5B base parameter count limits its reasoning depth on complex multi-hop questions relative to frontier models. While runtime inference supports up to 4,096 tokens, training sequence length was constrained to 2,048 tokens, requiring chunked retrieval for longer documents. The reported NLP quality scores are design-informed theoretical estimates extrapolating from LoRA scaling literature rather than formal empirical benchmark sweeps; comprehensive empirical validation of task accuracy across standardized benchmarks remains as future work. Planned next steps include longer context windows, on-disk FAISS for multi-document corpora exceeding available RAM, and RLHF fine-tuning for improved response factuality. While on-device latency, memory consumption, and generation throughput have been empirically measured and validated on the reference hardware, formal accuracy benchmark

suites remain a priority for subsequent work.

## 8 Conclusion

We have presented **MNIME**, a fully local, offline multimodal document processing system that integrates the fine-tuned MNIME-Core-V5 model with a high-throughput C-accelerated PDF engine and a premium PyQt6 desktop interface. MNIME demonstrates that competitive document AI—including semantic search, cross-referencing, and conversational RAG—is achievable entirely on-device on consumer hardware, without any external API calls.

Theoretical performance projections suggest that MNIME-Core-V5 should achieve approximately 88–90% of GPT-4o-mini’s NLP quality on document-specific tasks after LoRA fine-tuning, while the O(1) carousel architecture is designed to maintain 60 FPS responsiveness with up to 5,000 loaded files. On-device empirical testing has validated sub-100 ms first-token latency and 56 tokens/s generation on the reference AMD Ryzen 9 / RTX 3060 workstation, while comprehensive empirical benchmarking of end-to-end task accuracy is planned as a follow-up. We release MNIME as open-source software to support the growing community of researchers and practitioners who require private, high-quality document AI at the edge.

## 9 Limitations: Model Agency and Training Revisions

While MNIME-Core-V5 demonstrates strong empirical performance on deterministic document synthesis and RAG, it fundamentally failed to exhibit spontaneous agency or “come alive” as a proactive, autonomous entity. Despite the philosophical discourse embedded in the V4 and V5 datasets, the model remains strictly reactive, producing responses only when explicitly prompted by a user or system event. Additionally, empirical testing revealed an increased propensity for logical hallucination, which is likely a direct consequence of “super-charging” the V5 ethics and anti-bias pipeline at the expense of deterministic precision. Overcoming these limitations requires moving beyond static instruction-tuning pairs toward active agentic frameworks, reinforcement learning from environment feedback, and continuous state-memory architectures. Consequently, the core training pipeline is slated for a fundamental revision fol-

lowing further architectural planning. Future training iterations will explicitly include the original 1.5B model in each phase as a base profile class to ensure its foundational document processing edge and logical focus are not obscured by subsequent value alignments.

## 10 Future Work: Deep Semantic Integration

The current architecture supports conversational RAG alongside high-performance document manipulation. The next phase of development focuses on deep semantic integration, fusing the MNIME-Core engine directly into the core processing tools. Proposed advancements include **Semantic Document Splitting** (using NLP to detect topic boundaries and classify documents), **Enhanced “Smart” Bookmarking** (upgrading heuristic outline generation to semantic summarization and entity extraction), **NLP-Augmented Viewing** (embedding summarization and semantic bounding-box highlighting directly into the reading canvas), and **Intelligent Merging** (synthesizing auto-generated executive summary cover pages during document concatenation).

## Ethics Statement

MNIME is designed with user privacy as a first-order constraint. No user data, document content, or usage telemetry is transmitted externally. The fine-tuning data was drawn exclusively from publicly available, permissively licensed document corpora. Performance figures for cross-referencing and QA accuracy are theoretical projections grounded in published LoRA scaling benchmarks, while model latency, memory footprint, context ingestion, and generation throughput have been empirically measured on the reference workstation (AMD Ryzen 9 7950X, RTX 3060 Ventus 12 GB). The system is intended to augment, not replace, human judgment in document review.

## Statement on AI Use

The authors used AI writing assistants to improve clarity and LaTeX formatting in preparation of this manuscript.

## References

- Yuvanesh Anand, Zach Nussbaum, Brandon Duderstadt, Benjamin Schmidt, and Andriy Mulyar. 2023. GPT4All: An ecosystem of open-source on-edge large language models. <https://gpt4all.io>.
- Artifex Software. 2024. PyMuPDF: Python bindings for MuPDF. <https://pymupdf.readthedocs.io>.
- Adrian Berry. 1983. *The Super-Intelligent Machine*. Jonathan Cape, London.
- Harrison Chase. 2022. LangChain: Building applications with LLMs through composability. <https://github.com/langchain-ai/langchain>.
- Gemini Team, Rohan Anil, Sebastian Borgeaud, et al. 2025. Gemini: A family of highly capable multi-modal models. *Preprint, arXiv:2312.11805*.
- Georgi Gerganov. 2023. llama.cpp: Llm inference in C/C++. <https://github.com/ggerganov/llama.cpp>.
- Edward Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2021. LoRA: Low-rank adaptation of large language models. *Preprint, arXiv:2106.09685*.
- Yupan Huang, Tengchao Lv, Lei Cui, Yutong Lu, and Furu Wei. 2022. LayoutLMv3: Pre-training for document AI with unified text and image masking. In *Proceedings of the 30th ACM International Conference on Multimedia*, pages 4083–4091.
- Jeff Johnson, Matthijs Douze, and Hervé Jégou. 2019. Billion-scale similarity search with GPUs. In *IEEE Transactions on Big Data*, volume 7, pages 535–547. IEEE.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33:9459–9474.
- Jerry Liu. 2022. LlamaIndex: Data framework for LLM applications. [https://github.com/jerryliu/llama\\_index](https://github.com/jerryliu/llama_index).
- OpenAI. 2024. GPT-4 technical report. *Preprint, arXiv:2303.08774*.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. 2022. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744.
- Qwen Team. 2025. Qwen2.5: A party of foundation language models. *Preprint, arXiv:2412.15115*.
- Jason Wei, Maarten Bosma, Vincent Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M Dai, and Quoc V Le. 2022. Finetuned language models are zero-shot learners. *Preprint, arXiv:2109.01652*.

## A System Requirements and Installation

MNIME requires Python 3.10+ and has been tested on Windows 10/11. Two installation channels are available:

1. **Pre-built installer:** MNIME\_Setup.exe (Inno Setup wizard) or MNIME\_installer.exe (custom animated PyQt6 installer), available in the `installer/` directory.
2. **Source:** Clone the repository, download MNIME-Core-V5-Q4\_K\_M.gguf from Hugging Face ([https://huggingface.co/KyleDeanML/MNIME-Core-V5-Q4\\_K\\_M](https://huggingface.co/KyleDeanML/MNIME-Core-V5-Q4_K_M)), place it in `models/`, then run `setup.bat` followed by `run.bat`.

GPU acceleration is auto-detected at runtime via llama-cpp-python's CUDA backend. NVIDIA GPU layers are set to -1 by default, offloading as many transformer layers as will fit in available VRAM.

## B Key Performance Optimizations

**C-Accelerated Processing.** Merge, extraction, and compression operations use native C-level PyMuPDF routines, achieving up to 50× throughput improvement over pure-Python PDF libraries.

### Chunked Batch & NLP-Assisted Merging.

Single-pass document merging of thousands of files typically triggers process memory limits and C-level memory aborts in native libraries. MNIME introduces a two-tier chunked segment pipeline (batch size of 250) that caps active RAM consumption at < 200 MB regardless of total page count. During final assembly, rather than loading all segments simultaneously into memory, chunks are sequentially streamed to disk via PyMuPDF's incremental save (`saveIncr`), maintaining near-zero active footprint. Concurrently, the engine integrates NLP Engine during file ingestion to automatically build hierarchical TOC outlines and smart section bookmarks without requiring separate post-processing steps.

**O(1) Carousel Operations.** Drag-and-drop reordering and card removal use surgical layout-index shifts rather than rebuilding the widget tree, giving constant-time performance regardless of file count.

**In-Memory Caching.** Thumbnails and vector SVG icons are rasterized and pre-scaled once per session, eliminating CPU resampling during scroll and hover events.

**Non-Blocking Workers.** All file processing and NLP inference run in dedicated QThread workers. The main thread handles only UI events, guaranteeing frame-rate consistency.

**Dynamic NLP Memory.** The GGUF model and FAISS index are completely deallocated when NLP is toggled off or the application closes, preventing idle memory consumption.

**System Resource Purger.** On application exit, `core/system_cleaner.py` explicitly unloads the llama.cpp model (releasing its GPU allocation), drops the in-memory global vector store, runs garbage collection, and removes leftover temporary artifacts (`*.tmp`, `*.tmp.pdf`) from project directories, excluding environment, model, and training folders.

**V6 “Wild Tier” Memory Diet.** A 64KB C# managed system tray stub acts as the primary lifecycle manager. Combined with aggressive EcoQoS working set trimming, zero-dependency FAISS extraction, and view-port clipping, recent empirical telemetry during minimized idle operation demonstrated a working set of 136.0 MB (62.4 MB private) with a peak of 174.3 MB, operating on 2 threads and 1,183

handles, while consuming negligible CPU (0.001% or 0.02 s over a 60-second window). When completely hidden for 10 minutes, the Python process fully terminates while persisting the file queue to `session.json`, releasing 100% of computational resources until restored by the tray stub.

**OS Integration.** Deep Windows integration (`core/windows_integration.py`) manages taskbar progress states and native UI extensions, directly binding OS-level feedback to internal processing pipelines.